

114-1 TAICA Natural Language Processing: Term Project

Checkpoint 1 and 2 of WattBot 2025

Team name: Attention Please

November 29, 2025

YouTube Video Link:

<https://youtu.be/YavBekCeogE>

Outline

- 1 Section 1: Competition Overview
- 2 Section 2: Dataset Specification
- 3 Section 3: Challenges and Problems
- 4 Section 4: Kaggle Competitor Solutions Review: KohakuRAG Pipeline
- 5 Section 5: Our Baseline Method: BM25
- 6 Section 6: Our HERO Pipeline (Hybrid Evidence Retrieval Optimization)
- 7 Section 7: Conclusion and Future Work

Section 1: Competition Overview

Context & Problem

- AI training consumes massive energy, but data is scattered in unstructured reports.
- **Gap:** Lack of transparent, evidence-based sustainability estimation.

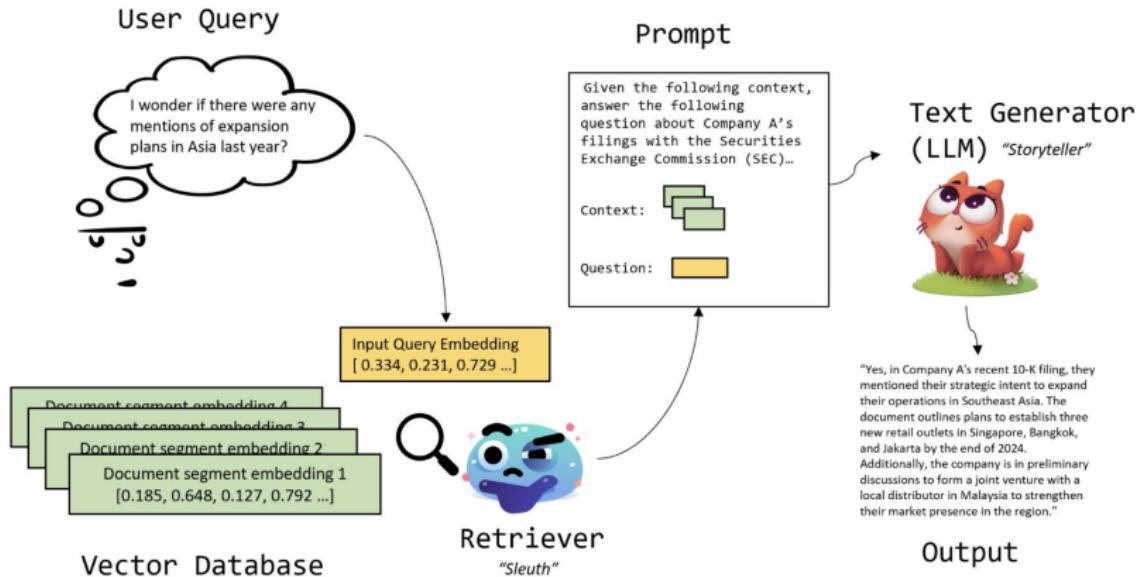
Goal: Verifiable RAG System

- Extract sustainability insights from 30+ academic papers.
- **Required Outputs:**
 - Value & Unit: Normalized numeric data (e.g., 1.2 MWh → 1200 kWh).
 - Ref_ID: Precise citations for traceability.

Evaluation Metrics

- **75% Accuracy:** Numeric values within $\pm 0.1\%$ tolerance.
- **15% Evidence:** Jaccard overlap of citation IDs.
- **10% Safety:** Correctly identify unanswerable queries (`is_blank`).

What is retrieval augmented generation (RAG)?



Section 2: Dataset Specification

1. Document Corpus (30+ PDFs)

- **Sources:** Academic papers (arXiv) & Industry reports (Amazon).
- **Challenge:** High heterogeneity (multi-column text, footnotes).
- **Complex Tables:** Spanning multiple pages, merged cells, borderless.
- **Hidden Data:** Key values embedded in charts without text descriptions.

2. Q&A Task Types

- **Numeric:** Extract specific values (e.g., "CO2 emissions of BERT").
- **Boolean:** Verify facts (e.g., "Did GPT-3 use more water than BLOOM?").
- **Terminology:** Define specific metrics (e.g., "Carbon Intensity").
- **Unanswerable:** Detect lack of evidence (hallucination test).

Section 3: Challenges and Problems - Data Extraction

1. The "Table Structure" Problem

- **Issue:** Standard parsers treat PDFs as linear text streams.
- **Impact:**
 - Row/Column misalignment in complex tables.
 - Values separated from headers (e.g., "1438" vs "lbs").
 - Cross-page tables broken by standard chunking.

2. "Invisible" Visual Data

- **Issue:** Critical trends often exist *only* in charts.
- **Impact:** Text-only pipelines miss this data completely (0% recall).

Section 3: Challenges and Problems - Reasoning Logic

3. Unit Ambiguity & Normalization

- **Issue:** Mixed units (kWh vs MWh, lbs vs kg) across papers.
- **Risk:** LLMs struggle with precise arithmetic conversion.
- **Constraint:** Score implies $\pm 0.1\%$ tolerance; approximation fails.

4. Hallucination Risks

- **Issue:** RAG tends to answer even irrelevant questions (e.g., "Elephant weight").
- **Penalty:** Severe score reduction for false positives.
- **Need:** A strict "Refusal Mechanism" for unanswerable queries.

Section 4: Case Study - KohakuRAG Pipeline (Indexing)

1. Hierarchical Indexing Structure

- **Tree-Based Parsing:** Documents are parsed into a semantic hierarchy:

Doc → Section → Paragraph → Sentence

- **Leaf Nodes (Sentences):**

- Embedded using `jina-embeddings-v3` (1024-dim).
- Used for high-granularity vector matching.

- **Parent Nodes (Paragraphs):**

- Inherit averaged vectors from children.
- Stored alongside leaves to provide **Full Context** upon retrieval.

2. Single-File Datastore (Key Differentiator)

- **Tech Stack:** Built on **SQLite + sqlite-vec** (via KohakuVault).
- **Benefit:** No complex vector DB (like Faiss/Milvus) required. Entire index is a single .db file, ensuring reproducibility and easy deployment.

Section 4: Case Study - KohakuRAG Pipeline (Retrieval)

3. Smart Context Expansion Strategy

- Leaf-to-Root Retrieval:

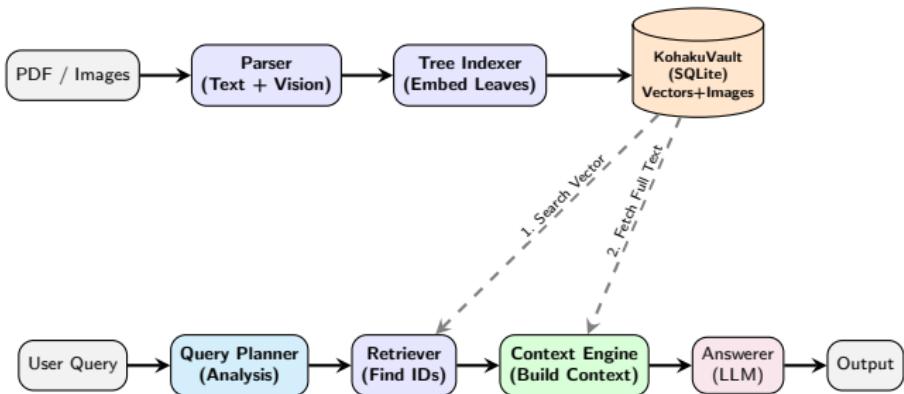
- ① Match: Query matches a specific Sentence (Leaf) via Vector Search.
- ② Expand: System automatically retrieves the Parent Paragraph and Section Metadata.
- ③ Result: The LLM sees the surrounding context, not just the isolated sentence.

4. Robust Production Features

- Rate-Limit Resilience: Implemented Exponential Backoff and Semaphore-based concurrency to handle OpenAI API limits automatically.
- Multimodal Integration:

- Uses Vision Models (e.g., Qwen-VL) to generate captions for figures.
- Captions are indexed alongside text, enabling retrieval of visual data.

Section 4: Case Study - KohakuRAG Pipeline (Architecture)



Pipeline Logic:

- **Planner:** Decides *how* to search (e.g., keyword extraction).
- **Fetch Full Text:** Vector search only finds IDs. We must query the DB again to retrieve the actual paragraph text and parent context.

Section 4: Case Study - Key Innovations

1. Tree-Based Context Expansion

- **Innovation:** Matches query against detailed sentences (Leaf) but automatically retrieves the full surrounding paragraph (Parent).
- **Benefit:** Solves "Context Fragmentation." The LLM gets coherent stories, not just isolated fragments.

2. Multimodal Layout Parsing

- **Technique:** Inserts tags like [Image page=1 idx=1] and stores generated captions (via Qwen-VL) in the same DB.
- **Benefit:** Preserves document flow and allows retrieving visual evidence (charts/figures) alongside text.

3. Production-Grade Engineering

- **Single-File Architecture:** Built on **SQLite** (KohakuVault). No external dependencies (like Pinecone/Docker), ensuring 100% reproducibility.
- **Async I/O:** Full asyncio pipeline handles API rate limits (Backoff) automatically, maximizing throughput.

What can we learn from the case study?

1. Structure > Chunking

- **Lesson:** Do not treat PDFs as flat text strings. Preserving document hierarchy (Doc → Section → Leaf) allows precise targeting.

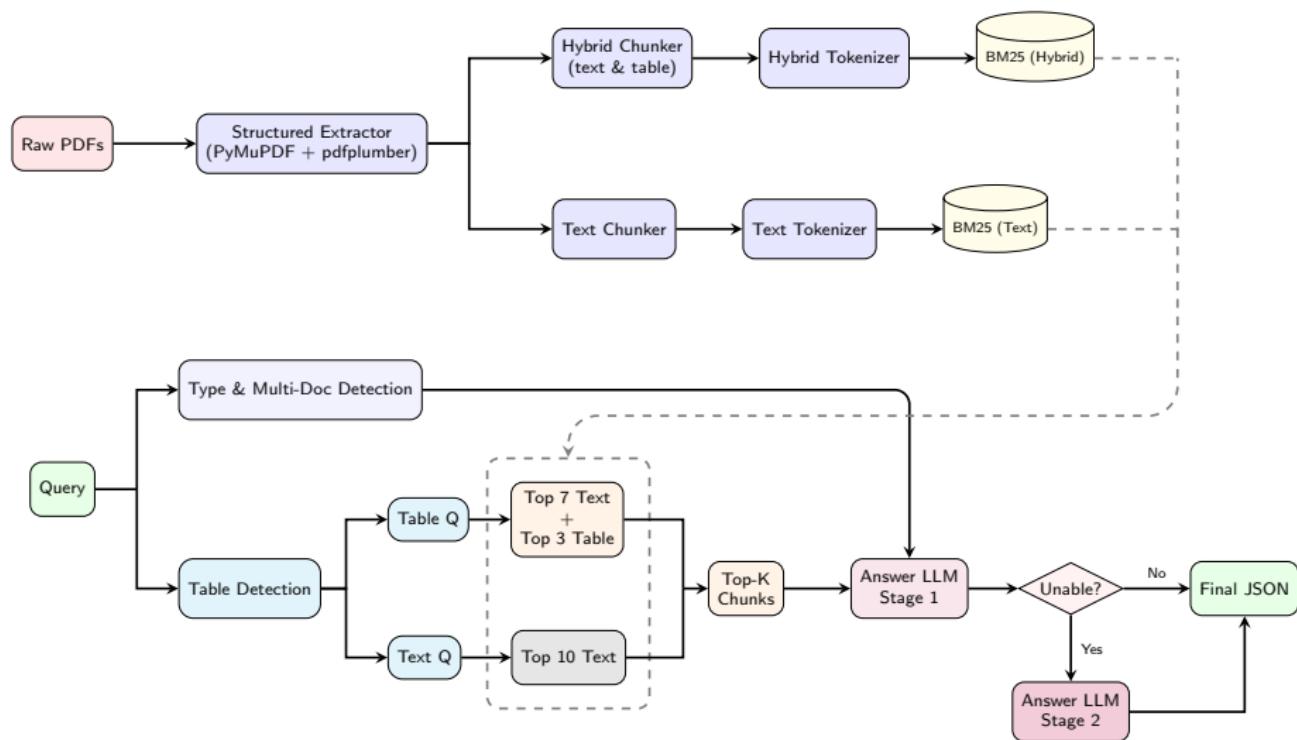
2. Engineering Simplicity

- **No Vector DB:** Using **SQLite** as a single-file datastore reduces complexity and ensures reproducibility.
- **Resilience:** Handling API rate limits with **Backoff** logic is crucial for batch processing without crashing.

3. Visuals are Data, Not Noise

- **Lesson:** Critical energy metrics are often hidden in charts. We must process images (via Vision Models/OCR), not just filter them out.

Section 5: Our Baseline Method — BM25 Sparse Retrieval



Section 5: Our Baseline Method — Data Ingestion Pipeline

1. Document Parsing Strategy

- **Tools:** PyMuPDF and pdfplumber (Standard CPU extraction).
- **Objective:** Generated two datasets to isolate the impact of "Structure Noise".

Experimental Setup: Two Dataset Versions

A. Hybrid Set (2,608 chunks)

- Text + Raw Markdown Tables
- (*Hypothesis: More info is better*)

B. Text-Only Set (1,778 chunks)

- Pure text, tables removed
- (*Hypothesis: Less noise is better*)

2. Enhanced Tokenization Logic

Standard tokenizers split numbers, losing numeric precision. We fixed this with a custom Regex to preserve numerical entities.

```
def tokenize(text: str) -> List[str]:  
    text = text.lower()  
    # Pattern: Preserve floats, percentages, or any word character  
    tokens = re.findall(r'\d+\.\?\d*%?|\w+', text)  
    return tokens
```

Impact: "H100 usage: 85.5%" → ['h100', 'usage', '85.5%'] (**Value Preserved**)

Section 5: Our Baseline Method — Retrieval Strategy

1. Query Routing Mechanism

- Implemented a keyword-based router to detect table-dependent questions.
- **Keywords:** table, how many, compare, versus, GB.

2. Weighting Logic (The Experiment)

Query Type	Retrieval Source	Weight
Text Query	Text-Only Index (Top 10)	1.0
Table Query	Text-Only Index (Top 7)	1.0
	Hybrid Index (Top 3)	0.8

We deliberately applied a **0.8 weight penalty** to table chunks because raw Markdown formatting (e.g., |---|) acts as noise for BM25.

Section 5: Our Baseline Method — Retrieval Performance Analysis

Comprehensive Metrics

Pipeline	Recall@1	Recall@3	Recall@5	Recall@10	MRR
Hybrid (Noise)	79.49%	89.74%	92.31%	92.31%	0.8526
Text-Only (Clean)	87.18%	92.31%	92.31%	94.87%	0.8974

nDCG Ranking Quality

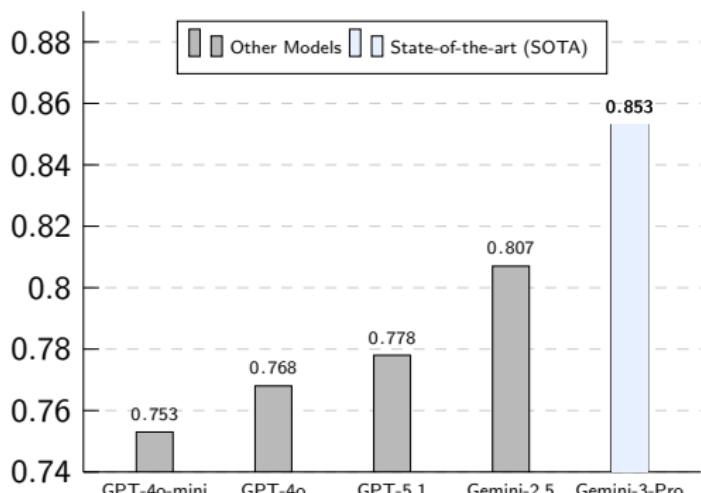
Pipeline	nDCG@1	nDCG@3	nDCG@5	nDCG@10
Hybrid (Noise)	0.7949	0.8497	0.8564	0.8617
Text-Only (Clean)	0.8718	0.8859	0.8859	0.8950

Key Insight: The "Structure Noise" Paradox

- **Text-Only Wins All Metrics:** The pipeline is superior in both coverage & ranking quality.
- **Diagnosis:** Raw table formatting acts as noise, **diluting** the Term Frequency (TF) score of actual keywords in the Hybrid index.
- **Conclusion:** BM25 cannot effectively leverage table structure, BM25 is not enough.

Section 5: Our Baseline Method — End-to-End Leaderboard

Model Performance (282 Test Queries)



Key Insight: The Retrieval Ceiling

- **LLM Power:** Gemini 3 Pro pushed the baseline score to **0.853**.
- **Bottleneck Remains:** Even with the strongest model, we are limited by the quality of BM25 retrieval (missing visual/table data).

Question: Why Focus on Retrieval?

- To break the ceiling further, we must improve **context quality** (e.g., extracting charts), not just switch LLMs.
- **Efficiency:** Baseline BM25 is cheap ($\sim \$0.003/\text{query}$) but fundamentally noisy.

Section 5: Our Baseline Method — Limitations and Future Direction

Despite achieving **0.853**, the Baseline hits three fundamental limits:

① Semantic Gap (Recall Issue)

BM25 fails on synonyms (e.g., "Carbon footprint" \neq "Emission"). We miss documents that use different terminology.

② The Table Dilemma (Data Loss)

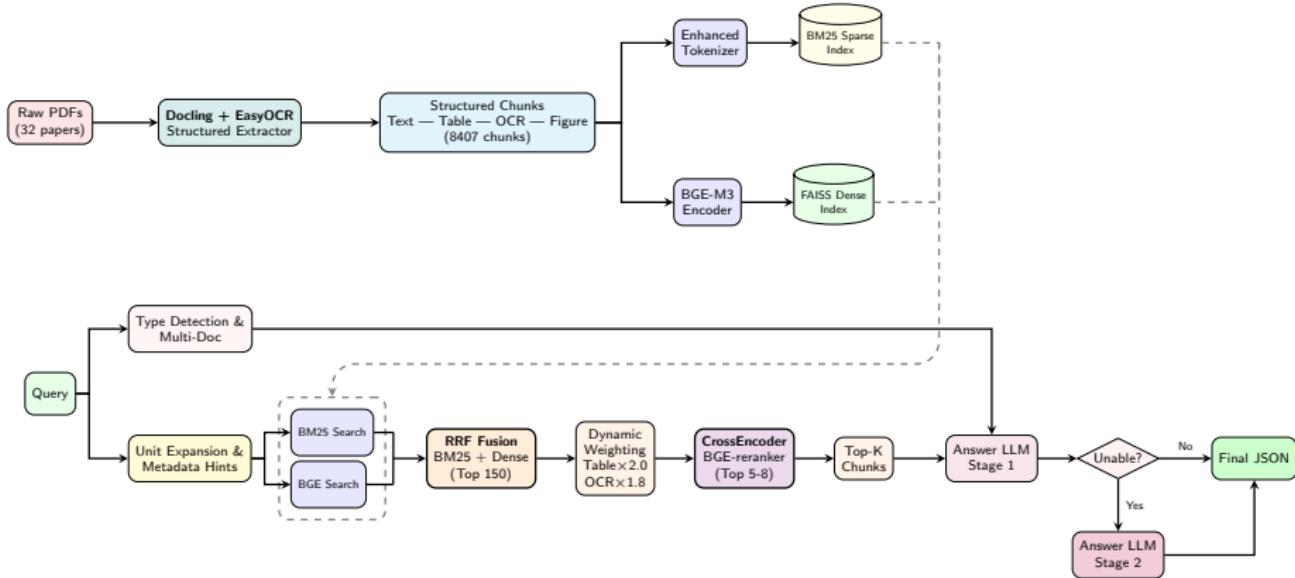
We suppressed tables (0.8 penalty) to reduce noise, but tables contain **90%** of the numerical answers. We are actively filtering out the truth.

③ Precision Limit (Hallucination Risk)

Without a Reranker, Top-K results contain irrelevant chunks, which confuses the LLM and leads to hallucinations.

Solution: Introducing the [HERO Pipeline!](#)

Section 6: HERO Pipeline (Hybrid Evidence Retrieval Optimization)



Section 6: HERO Pipeline - Visual-Aware Ingestion

1. Hybrid Parsing Pipeline (Docling + EasyOCR)

- **Problem:** Standard parsers ignore data embedded in charts/images.
- **Solution:**
 - Used Docling to classify chunks into Text, Table, and Picture.
 - **OCR Injection:** Automatically detected figures and applied EasyOCR (w/ OpenCV preprocessing) to extract numerical data from charts.
 - **Result:** Converted "invisible" pixels into searchable text chunks.

2. Enhanced Tokenization

- **Logic:** Custom tokenizer preserves crucial non-standard tokens:
- **Rules:** Keeps floating point numbers, years (202x), and specific technical terms, preventing BM25 from stripping vital numeric contexts.

Section 6: HERO Pipeline - Heuristic-Weighted RRF

3. Adaptive Reciprocal Rank Fusion (RRF)

- **Concept:** Combine BM25 (Sparse) and BGE-Large (Dense) with *content-aware weights*.
- **Weighted Formula:**

$$Score(d) = \sum \frac{1}{k + rank(d)} \times W_{type}$$

- **Heuristic Weights (W_{type}):**

- **Tables** ($\times 2.0$): High information density for metrics.
- **OCR/Images** ($\times 1.8$): Unique data often found only in charts.
- **Metadata Hints** ($\times 2.0$): Matches on Author/Year in query.

4. Semantic Unit Expansion

- **Technique:** Query Augmentation via UNIT_KEYWORDS.
- **Example:** Query "tCO2e" → expands to "carbon emissions co2 greenhouse footprint".
- **Benefit:** Bridges the gap between abbreviated metrics and natural language text.

Section 6: HERO Pipeline - Two-Stage Adaptive Generation

5. Dynamic Prompting Taxonomy

- **Classifier:** Regex-based router detects question type (e.g., Table_Extraction, Score_Extraction, Model_Specific).
- **Action:** Injects specialized System Prompts (e.g., "Search row-by-row" for tables).

6. Two-Stage Fallback Mechanism

- **Stage 1 (Precision):** Strict mode using Top-3 chunks. Low temperature to prevent hallucination.
- **Stage 2 (Recall):** Triggered if Stage 1 returns "Unable to Answer".
 - Expands to Top-10 chunks.
 - Switches to "Aggressive Mode": Relaxed constraints to capture partial matches or indirect evidence.

Section 7: Conclusion and Future Work - Retrieval Performance

Metrics Analysis

Metric	Baseline	HERO
Recall@K (Coverage)		
Recall@1	79.49%	80.49%
Recall@3	89.74%	89.80%
Recall@5	92.31%	92.68%
Recall@10	92.31%	95.12%
nDCG@K (Ranking Quality)		
nDCG@1	0.7949	0.8049
nDCG@3	0.8497	0.8502
nDCG@5	0.8564	0.8584
nDCG@10	0.8617	0.8665
Overall Accuracy		
MRR	0.8526	0.8560

Key Insights:

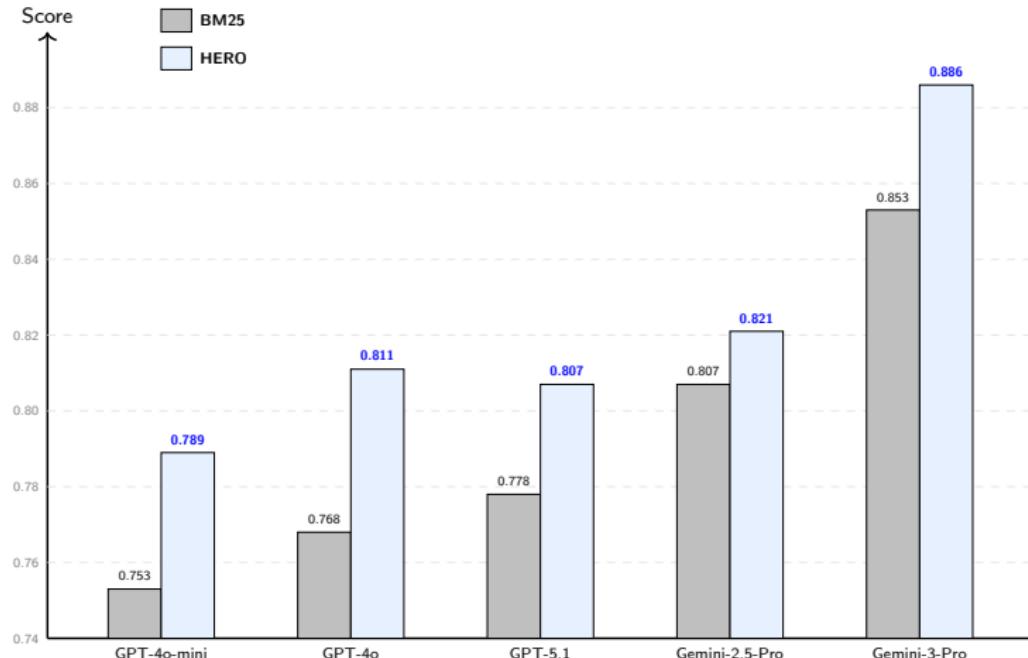
- **Precision Lift (Recall@1): +1% gain.** The Reranker successfully filtered the Baseline's "structure noise" and boosted ranking confidence.
- **Coverage Lift (Recall@10):** HERO reached 95.12% coverage, proving that Dense Retrieval effectively found semantic matches that BM25 missed.
- **Ranking Quality (nDCG):** Overall ranking improved by +1.7%. This stabilization proves our Dynamic Weighting ($\times 2.0$ for Tables) works.

Key Takeaway:

HERO approach converts the Baseline's weakest point (noisy multimodal data) into its greatest strength, ensuring the LLM receives optimal context.

Section 7: Conclusion and Future Work - Pipeline & Model Benchmarking

Comparison: Baseline method vs. HERO



Section 7: Conclusion and Future Work - Achievements

1. Zero-to-One: Visual Data Unlocked

- **Challenge:** Critical energy metrics (e.g., GPU trends) are often trapped in uncopyable charts.
- **Achievement:** Integrated EasyOCR into the ingestion pipeline, successfully converting pixel-based data into searchable text chunks.
- **Impact:** Expanded information coverage from text-only to **Multimodal (Text + Vision)**.

2. Structure-Aware Precision

- **Challenge:** Standard RAG treats high-density tables as normal text, diluting numeric accuracy.
- **Achievement:** Implemented **Heuristic-Weighted RRF** (Table $\times 2.0$, OCR $\times 1.8$).
- **Impact:** The weighted retrieval directly contributed to the **0.821** score by surfacing correct numeric answers over general descriptions.

Section 7: Conclusion and Future Work - Key Insights

Insight 1: Unit Semantic Gap

Problem:

Vector models struggle with
"tCO2e" ≠ "Carbon Footprint"

Solution:

Hard-coded Query Expansion

- Map MWh → *Megawatt Hour*
- Map tCO2e → *Carbon Emission*

Semantic search alone is insufficient for technical metrics.

Insight 2: Safety vs. Recall

Problem:

Strict prompt → High precision
But excessive "*Unable to Answer*"

Solution:

Two-Stage Fallback

- Stage 1:** Strict rules
(Top-3, Low temp)
- Stage 2:** Aggressive mode
(Top-10, Relaxed constraints)

Balance precision with coverage.

WattBot 2025

Submit Prediction

...

Overview Data Code Models Discussion Leaderboard Rules Team Submissions

#	Team	Members	Score	Entries	Last	Join
1	Attention Please		0.886	42	2d	

Kaggle Leaderboard: Our HERO pipeline achieved 0.886! Top 1

[Reminder] Scoring Criteria for Checkpoint 2 (For Reference Only)

Based on the evaluation criteria provided by the TAs.

Criteria	Score	Description
Creativity	0	Off-topic or shows no relevant creative contribution.
	1	Minimal effort in idea generation; no sign of originality.
	2	Mostly standard and predictable; lacks noticeable creative elements.
	3	Reasonably solid but follows conventional patterns; limited originality or new perspectives.
	4	Shows some originality or thoughtful variation in approach; includes at least one creative or non-obvious idea.
	5	Demonstrates highly original ideas or approaches; shows clear insight and innovative thinking beyond standard solutions.
Analysis Completeness	0	Incomplete, irrelevant, or incorrect analysis.
	1	Minimal analysis with little evidence of understanding.
	2	Partial or superficial analysis; several important aspects are missing or unclear.
	3	Addresses the main parts of the topic but misses some relevant points or details.
	4	Mostly complete and accurate analysis with minor gaps or small logical weaknesses.
	5	Thorough, accurate, and logically consistent analysis; all key aspects of the topic are addressed and well-supported.
Presentation Quality	0	No meaningful presentation delivered.
	1	Disorganized and unclear; audience has difficulty understanding the content.
	2	Hard to follow due to poor structure or unclear explanation.
	3	Understandable overall, though with some disorganization or inconsistent delivery.
	4	Clear and organized presentation with good pacing and minor issues in clarity or engagement.
	5	Extremely clear, well-structured, and engaging; visuals, timing, and delivery are polished and professional.

[Action Required] Self-Assessment and Scoring Alignment

Ready to score our presentation? We need your support!

Creativity

4-5

- ✓ Weighted RRF
Table ×2.0 boost
- ✓ Two-Stage LLM
Smart fallback
- ✓ Beyond Baseline
Learned from
competitors
- ✓ HERO Pipeline
No.1 in the Kaggle
Leader-board.

Analysis

4-5

- ✓ 7 Sections
Complete coverage
- ✓ 9 Metrics
Recall, nDCG, MRR
- ✓ 5 Models
Full benchmark
- ✓ Ablation Study
Text vs Hybrid

Presentation

4-5

- ✓ 3 Diagrams
TikZ architecture
- ✓ Clear Flow
Logical structure
- ✓ Unified Theme
Watt Blue design
- ✓ Data Visuals
6 tables + 1 chart

Self-Assessment Range: 12–15 points. Thank you for your attention. Much appreciated!